

labassignment8bn

March 24, 2025

1 Lab Assignment 8: Data Management Using pandas, Part 1

1.1 DS 6001: Practice and Application of Data Science

1.1.1 Instructions

Please answer the following questions as completely as possible using text, code, and the results of code as needed. Format your answers in a Jupyter notebook. To receive full credit, make sure you address every part of the problem, and make sure your document is formatted in a clean and professional way.

In this lab, you will be working with the [2017 Workplace Health in America survey](#) which was conducted by the Centers for Disease Control and Prevention. According to the survey's [guidance document](#):

The Workplace Health in America (WHA) Survey gathered information from a cross-sectional, nationally representative sample of US worksites. The sample was drawn from the Dun & Bradstreet (D&B) database of all private and public employers in the United States with at least 10 employees. Like previous national surveys, the worksite served as the sampling unit rather than the companies or firms to which the worksites belonged. Worksites were selected using a stratified simple random sample (SRS) design, where the primary strata were ten multi-state regions defined by the Centers for Disease Control and Prevention (CDC), plus an additional stratum containing all hospital worksites.

The data contain over 300 features that report the industry and type of company where the respondents are employed, what kind of health insurance and other health programs are offered, and other characteristics of the workplaces including whether employees are allowed to work from home and the gender and age makeup of the workforce. The data are full of interesting information, but in order to make use of the data a great deal of data manipulation is required first.

1.2 Problem 0

Import the following libraries:

```
[1]: import numpy as np
import pandas as pd
import sidetable
import sqlite3
import warnings
warnings.filterwarnings('ignore')
```

1.3 Problem 1

The raw data are stored in an ASCII file on the 2017 Workplace Health in America survey [homepage](#). Load the raw data directly into Python without downloading the data onto your harddrive and display a dataframe with only the 14th, 28th, and 102nd rows of the data. [1 point]

```
[2]: whpps = pd.read_csv('https://www.cdc.gov/workplace-health-promotion/media/files/
↳2024/06/whpps_120717.csv', sep='~')
whpps.iloc[[14,28,104],:]
```

```
[2]:      OC1    OC3    HI1    HI2    HI3    HI4    HRA1    HRA1A    HRA1B    HRA1E    ...    WL3_05  \
14      7      2.0    2.0    1.0    2.0    1.0    1.0      3.0      2.0      2.0    ...      NaN
28      1      3.0    2.0    3.0    1.0    1.0    2.0     96.0     96.0     96.0    ...      NaN
104     7     97.0   97.0   96.0   97.0    1.0   97.0     96.0     96.0     96.0    ...      NaN

      E1_09  Suppquex      Id  Region  CDC_Region  Industry  Size  Varstrata  \
14      NaN        2.0  1539.0    2.0         4.0         7.0   5.0         0.0
28      NaN        2.0  2755.0    3.0         5.0         7.0   6.0         0.0
104     NaN        2.0 12982.0    1.0         1.0         7.0   8.0         0.0

      Finalwt_worksite,,,,
14      47.793940929,,,,
28      47.793940929,,,,
104     47.793940929,,,,

[3 rows x 301 columns]
```

1.4 Problem 2

The data contain 301 columns. Create a new variable in Python's memory to store a working version of the data. In the working version, delete all of the columns except for the following:

- **Industry:** 7 Industry Categories with NAICS codes
- **Size:** 8 Employee Size Categories
- **OC3** Is your organization for profit, non-profit, government?
- **HI1** In general, do you offer full, partial or no payment of premiums for personal health insurance for full-time employees?
- **HI2** Over the past 12 months, were full-time employees asked to pay a larger proportion, smaller proportion or the same proportion of personal health insurance premiums?
- **HI3:** Does your organization offer personal health insurance for your part-time employees?
- **CP1:** Are there health education programs, which focus on skill development and lifestyle behavior change along with information dissemination and awareness building?
- **WL6:** Allow employees to work from home?
- Every column that begins **WD**, expressing the percentage of employees that have certain characteristics at the firm

[1 point]

```
[3]: wd = [x for x in whpps.columns if x.startswith('WD')]
cols = ['Industry', 'Size', 'OC3', 'HI1', 'HI2', 'HI3', 'CP1', 'WL6'] + wd
whpps = whpps[cols]
whpps.head()
```

```
[3]:
```

	Industry	Size	OC3	HI1	HI2	HI3	CP1	WL6	WD1_1	WD1_2	WD2	WD3	\
0	7.0	7.0	3.0	2.0	1.0	2.0	1.0	1.0	25.0	20.0	85.0	60.0	
1	7.0	6.0	3.0	2.0	3.0	1.0	1.0	1.0	997.0	997.0	90.0	90.0	
2	7.0	8.0	3.0	1.0	3.0	1.0	1.0	1.0	35.0	4.0	997.0	997.0	
3	7.0	4.0	2.0	1.0	2.0	1.0	2.0	2.0	50.0	15.0	50.0	85.0	
4	7.0	4.0	3.0	1.0	3.0	1.0	1.0	1.0	50.0	40.0	60.0	60.0	

	WD4	WD5	WD6	WD7
0	40.0	15.0	0.0	22.0
1	997.0	997.0	0.0	997.0
2	40.0	15.0	997.0	997.0
3	75.0	0.0	0.0	997.0
4	40.0	30.0	0.0	28.0

1.5 Problem 3

The [codebook](#) for the WHA data contain short descriptions of the meaning of each of the columns in the data. Use these descriptions to decide on better and more intuitive names for the columns in the working version of the data, and rename the columns accordingly. [1 point]

```
[4]: whpps = whpps.rename({'OC3': 'org_status',
                           'HI1': 'premium_payment',
                           'HI2': 'premium_prop',
                           'HI3': 'part-time_insurance',
                           'CP1': 'health_education',
                           'WL6': 'telework',
                           'WD1_1': 'under_30',
                           'WD1_2': 'over_60',
                           'WD2': 'female',
                           'WD3': 'hourly_nonexempt',
                           'WD4': 'nondaytime',
                           'WD5': 'remote_offsite',
                           'WD6': 'union',
                           'WD7': 'turnover'}, axis=1)
whpps.head()
```

```
[4]:
```

	Industry	Size	org_status	premium_payment	premium_prop	\
0	7.0	7.0	3.0	2.0	1.0	
1	7.0	6.0	3.0	2.0	3.0	
2	7.0	8.0	3.0	1.0	3.0	
3	7.0	4.0	2.0	1.0	2.0	

4	7.0	4.0	3.0	1.0	3.0	
---	-----	-----	-----	-----	-----	--

	part-time_insurance	health_education	telework	under_30	over_60	female \
0	2.0	1.0	1.0	25.0	20.0	85.0
1	1.0	1.0	1.0	997.0	997.0	90.0
2	1.0	1.0	1.0	35.0	4.0	997.0
3	1.0	2.0	2.0	50.0	15.0	50.0
4	1.0	1.0	1.0	50.0	40.0	60.0

	hourly_nonexempt	nondaytime	remote_offsite	union	turnover
0	60.0	40.0	15.0	0.0	22.0
1	90.0	997.0	997.0	0.0	997.0
2	997.0	40.0	15.0	997.0	997.0
3	85.0	75.0	0.0	0.0	997.0
4	60.0	40.0	30.0	0.0	28.0

1.6 Problem 4

Using the codebook and this [dictionary of NAICS industrial codes](#), place descriptive labels on the categories of the industry column in the working data. [1 point]

```
[5]: industry_map = {1: 'Manufacturing & Construction',
                     2: 'Trade & Transportation',
                     3: 'Hospitality & Services',
                     4: 'Finance & Technology',
                     5: 'Education & Healthcare',
                     6: 'Public Administration',
                     7: 'Hospital worksites' }
whpps.Industry = whpps.Industry.map(industry_map)
whpps.Industry
```

```
[5]: 0      Hospital worksites
     1      Hospital worksites
     2      Hospital worksites
     3      Hospital worksites
     4      Hospital worksites
     ...
2838    Public Administration
2839    Public Administration
2840    Public Administration
2841    Public Administration
2842    Public Administration
Name: Industry, Length: 2843, dtype: object
```

1.7 Problem 5

Using the codebook, recode the “size” column to have three categories: “Small” for workplaces with fewer than 100 employees, “Medium” for workplaces with at least 100 but fewer than 500

employees, and “Large” for companies with at least 500 employees. [Note: Python dataframes have an attribute `.size` that reports the space the dataframe takes up in memory. Don’t confuse this attribute with the column named “Size” in the raw data.] [1 point]

```
[6]: size_map = {1: 'Small',
                2: 'Small',
                3: 'Small',
                4: 'Medium',
                5: 'Medium',
                6: "Large",
                7: "Large",
                8: "Large"}
whpps.Size = whpps.Size.map(size_map)
whpps.Size
```

```
[6]: 0      Large
     1      Large
     2      Large
     3    Medium
     4    Medium
     ...
    2838    Medium
    2839    Medium
    2840      Large
    2841      Large
    2842      Large
     Name: Size, Length: 2843, dtype: object
```

1.8 Problem 6

Use the codebook to write accurate and descriptive labels for each category for each categorical column in the working data. Then apply all of these labels to the data at once. Code “Legitimate Skip”, “Don’t know”, “Refused”, and “Blank” as missing values. [2 points]

```
[7]: replace_map = {'org_status': {1: 'For Profit', 2: 'For Profit', 3: 'Non-profit',
    ↪4: 'Government', 5: 'Government', 6: np.nan, 7: np.nan, 8: np.nan},
    ↪      'premium_payment': {1: 'Full', 2: 'Partial', 3: 'No', 97: np.nan,
    ↪98: np.nan, 99: np.nan},
    ↪      'premium_prop': {1: 'Full', 2: 'Partial', 3: 'No', 97: np.nan,
    ↪98: np.nan, 99: np.nan},
    ↪      'part-time_insurance': {1: 'Yes', 2: 'No', 97: np.nan, 98: np.
    ↪nan, 99: np.nan},
    ↪      'health_education': {1: 'Yes', 2: 'No', 97: np.nan, 98: np.nan},
    ↪      'telework': {1: 'Yes', 2: 'No', 97: np.nan, 98: np.nan, 99: np.
    ↪nan}}
whpps = whpps.replace(replace_map)
whpps.head()
```

```
[7]:
```

	Industry	Size	org_status	premium_payment	premium_prop	\
0	Hospital worksites	Large	Non-profit	Partial	Full	
1	Hospital worksites	Large	Non-profit	Partial	No	
2	Hospital worksites	Large	Non-profit	Full	No	
3	Hospital worksites	Medium	For Profit	Full	Partial	
4	Hospital worksites	Medium	Non-profit	Full	No	

	part-time_insurance	health_education	telework	under_30	over_60	female	\
0	No	Yes	Yes	25.0	20.0	85.0	
1	Yes	Yes	Yes	997.0	997.0	90.0	
2	Yes	Yes	Yes	35.0	4.0	997.0	
3	Yes	No	No	50.0	15.0	50.0	
4	Yes	Yes	Yes	50.0	40.0	60.0	

	hourly_nonexempt	nondaytime	remote_offsite	union	turnover
0	60.0	40.0	15.0	0.0	22.0
1	90.0	997.0	997.0	0.0	997.0
2	997.0	40.0	15.0	997.0	997.0
3	85.0	75.0	0.0	0.0	997.0
4	60.0	40.0	30.0	0.0	28.0

1.9 Problem 7

The features that measure the percent of the workforce with a particular characteristic use the codes 997, 998, and 999 to represent “Don’t know”, “Refusal”, and “Blank/Invalid” respectively. Replace these values with missing values for all of the percentage features at the same time. [1 point]

```
[8]: whpps = whpps.replace([997, 998, 999], np.nan)
      whpps.head()
```

```
[8]:
```

	Industry	Size	org_status	premium_payment	premium_prop	\
0	Hospital worksites	Large	Non-profit	Partial	Full	
1	Hospital worksites	Large	Non-profit	Partial	No	
2	Hospital worksites	Large	Non-profit	Full	No	
3	Hospital worksites	Medium	For Profit	Full	Partial	
4	Hospital worksites	Medium	Non-profit	Full	No	

	part-time_insurance	health_education	telework	under_30	over_60	female	\
0	No	Yes	Yes	25.0	20.0	85.0	
1	Yes	Yes	Yes	NaN	NaN	90.0	
2	Yes	Yes	Yes	35.0	4.0	NaN	
3	Yes	No	No	50.0	15.0	50.0	
4	Yes	Yes	Yes	50.0	40.0	60.0	

	hourly_nonexempt	nondaytime	remote_offsite	union	turnover
0	60.0	40.0	15.0	0.0	22.0
1	90.0	NaN	NaN	0.0	NaN

2	NaN	40.0	15.0	NaN	NaN
3	85.0	75.0	0.0	0.0	NaN
4	60.0	40.0	30.0	0.0	28.0

1.10 Problem 8

Sort the working data by industry in ascending alphabetical order. Within industry categories, sort the rows by size in ascending alphabetical order. Within groups with the same industry and size, sort by percent of the workforce that is under 30 in descending numeric order. [1 point]

```
[9]: whpps.sort_values(by=['Industry', 'Size', 'under_30'], ascending=[True, True,
↪False])
whpps.head()
```

```
[9]:
```

	Industry	Size	org_status	premium_payment	premium_prop	\
0	Hospital worksites	Large	Non-profit	Partial	Full	
1	Hospital worksites	Large	Non-profit	Partial	No	
2	Hospital worksites	Large	Non-profit	Full	No	
3	Hospital worksites	Medium	For Profit	Full	Partial	
4	Hospital worksites	Medium	Non-profit	Full	No	

	part-time_insurance	health_education	telework	under_30	over_60	female	\
0	No	Yes	Yes	25.0	20.0	85.0	
1	Yes	Yes	Yes	NaN	NaN	90.0	
2	Yes	Yes	Yes	35.0	4.0	NaN	
3	Yes	No	No	50.0	15.0	50.0	
4	Yes	Yes	Yes	50.0	40.0	60.0	

	hourly_nonexempt	nondaytime	remote_offsite	union	turnover
0	60.0	40.0	15.0	0.0	22.0
1	90.0	NaN	NaN	0.0	NaN
2	NaN	40.0	15.0	NaN	NaN
3	85.0	75.0	0.0	0.0	NaN
4	60.0	40.0	30.0	0.0	28.0

1.11 Problem 9

There is one row in the working data that has a NaN value for industry. Delete this row. Use a logical expression, and not the row number. [1 point]

```
[10]: null = np.where(whpps.Industry.isna())
null
whpps = whpps.drop(null[0])
```

1.12 Problem 10

Create a new feature named `gender_balance` that has three categories: “Mostly men” for workplaces with between 0% and 35% female employees, “Balanced” for workplaces with more than

35% and at most 65% female employees, and “Mostly women” for workplaces with more than 65% female employees. [1 point]

```
[11]: whpps = whpps.assign(gender_balance =
                                pd.cut(whpps.female,
                                        bins=[0, 35, 65, 100],
                                        labels=("Mostly men", "Balanced", "Mostly women")))
whpps.head()
```

```
[11]:
```

	Industry	Size	org_status	premium_payment	premium_prop	\
0	Hospital worksites	Large	Non-profit	Partial	Full	
1	Hospital worksites	Large	Non-profit	Partial	No	
2	Hospital worksites	Large	Non-profit	Full	No	
3	Hospital worksites	Medium	For Profit	Full	Partial	
4	Hospital worksites	Medium	Non-profit	Full	No	

	part-time_insurance	health_education	telework	under_30	over_60	female	\
0	No	Yes	Yes	25.0	20.0	85.0	
1	Yes	Yes	Yes	NaN	NaN	90.0	
2	Yes	Yes	Yes	35.0	4.0	NaN	
3	Yes	No	No	50.0	15.0	50.0	
4	Yes	Yes	Yes	50.0	40.0	60.0	

	hourly_nonexempt	nondaytime	remote_offsite	union	turnover	\
0	60.0	40.0	15.0	0.0	22.0	
1	90.0	NaN	NaN	0.0	NaN	
2	NaN	40.0	15.0	NaN	NaN	
3	85.0	75.0	0.0	0.0	NaN	
4	60.0	40.0	30.0	0.0	28.0	

	gender_balance
0	Mostly women
1	Mostly women
2	NaN
3	Balanced
4	Balanced

1.13 Problem 11

Change the data type of all categorical features in the working data from “object” to “category”. [1 point]

```
[12]: cats = ['Industry', 'Size', 'org_status', 'premium_payment', 'premium_prop',
              'part-time_insurance', 'health_education', 'telework', 'gender_balance']
whpps[cats] = whpps[cats].astype('category')
whpps.dtypes
```



```
[12]: Industry          category
      Size              category
      org_status        category
      premium_payment   category
      premium_prop      category
      part-time_insurance category
      health_education   category
      telework           category
      under_30           float64
      over_60            float64
      female             float64
      hourly_nonexempt   float64
      nondaytime         float64
      remote_offsite     float64
      union              float64
      turnover           float64
      gender_balance     category
      dtype: object
```

1.14 Problem 12

Filter the data to only those rows that represent small workplaces that allow employees to work from home. Then report how many of these workplaces offer full insurance, partial insurance, and no insurance. Use a function that reports the percent, cumulative count, and cumulative percent in addition to the counts. [1 point]

```
[13]: whpps.query('Size == "Small" & telework == "Yes"').stb.freq(['premium_payment'])
```

```
[13]: premium_payment  count    percent  cumulative_count  cumulative_percent
0             Full      324  46.285714             324          46.285714
1             Partial    310  44.285714             634          90.571429
2              No       66   9.428571             700         100.000000
```

1.15 Problem 13

Anything that can be done in SQL can be done with **pandas**. The next several questions ask you to write **pandas** code to match a given SQL query. But to check that the SQL query and **pandas** code yield the same result, create a new database using the **sqlite3** package and input the cleaned WHA data as a table in this database. (See module 6 for a discussion of SQLite in Python.) [1 point]

```
[ ]: whadb = sqlite3.connect('whadb.db')
     whpps.to_sql('whpps', whadb, index=False, chunksize=1000, if_exists = 'replace')
```

```
[ ]: 2842
```

1.16 Problem 14

Write **pandas** code that replicates the output of the following SQL code:

```
SELECT size, type, premiums AS insurance, percent_female FROM whpps
WHERE industry = 'Hospitals' AND premium_change='Smaller'
ORDER BY percent_female DESC;
```

For each of these queries, your feature names might be different from the ones listed in the query, depending on the names you chose in problem 3. [2 points]

```
[15]: myquery14 = '''
SELECT Size, org_status, premium_payment AS insurance, female FROM whpps
WHERE industry = 'Hospital worksites' AND premium_prop='Partial'
ORDER BY female DESC
'''
pd.read_sql_query(myquery14, whadb)
```

```
[15]:
```

	Size	org_status	insurance	female
0	Medium	Non-profit	Full	89.0
1	Large	Non-profit	Partial	80.0
2	Large	Non-profit	Partial	80.0
3	Small	Non-profit	Full	75.0
4	Medium	Non-profit	Partial	65.0
5	Medium	For Profit	Full	50.0
6	Medium	97.0	Partial	NaN
7	Medium	Non-profit	Partial	NaN
8	Medium	Non-profit	Full	NaN
9	Medium	Non-profit	Full	NaN
10	Large	Non-profit	Partial	NaN

```
[20]: whpps14 = whpps.query('Industry=="Hospital worksites" &
↳premium_prop=="Partial"').sort_values(by='female', ascending=False)
col14 = ['Size', 'org_status', 'premium_payment', 'female']
whpps14 = whpps14[col14]
whpps14 = whpps14.rename({'premium_payment': 'insurance'}, axis=1)
whpps14
```

```
[20]:
```

	Size	org_status	insurance	female
320	Medium	Non-profit	Full	89.0
187	Large	Non-profit	Partial	80.0
214	Large	Non-profit	Partial	80.0
229	Small	Non-profit	Full	75.0
191	Medium	Non-profit	Partial	65.0
3	Medium	For Profit	Full	50.0
11	Medium	97.0	Partial	NaN
48	Medium	Non-profit	Partial	NaN
51	Medium	Non-profit	Full	NaN
75	Medium	Non-profit	Full	NaN
97	Large	Non-profit	Partial	NaN

1.17 Problem 15

Write pandas code that replicates the output of the following SQL code:

```
SELECT industry,
       AVG(percent_female) as percent_female,
       AVG(percent_under30) as percent_under30,
       AVG(percent_over60) as percent_over60
FROM whpps
GROUP BY industry
ORDER BY percent_female DESC;
```

[2 points]

```
[30]: myquery15 = '''
SELECT industry,
       AVG(female) as percent_female,
       AVG(under_30) as percent_under30,
       AVG(over_60) as percent_over60
FROM whpps
GROUP BY industry
ORDER BY percent_female DESC
'''
pd.read_sql_query(myquery15, whadb)
```

```
[30]:
```

	Industry	percent_female	percent_under30	\
0	Education & Healthcare	80.657143	25.745665	
1	Hospital worksites	76.427027	27.213793	
2	Hospitality & Services	53.804416	38.566343	
3	Finance & Technology	50.632184	23.821752	
4	Public Administration	39.056738	21.015625	
5	Trade & Transportation	32.657258	29.108696	
6	Manufacturing & Construction	20.328605	22.257143	

	percent_over60
0	11.349570
1	16.489655
2	11.544872
3	12.465465
4	15.015385
5	12.584034
6	10.690355

```
[29]: whpps15 = whpps.groupby('Industry').agg({'female': 'mean', 'under_30': 'mean',
↪      'over_60': 'mean'}).sort_values('female', ascending=False)
whpps15 = whpps15.reset_index()
whpps15
```

```
[29]:
```

	Industry	female	under_30	over_60
0	Education & Healthcare	80.657143	25.745665	11.349570
1	Hospital worksites	76.427027	27.213793	16.489655
2	Hospitality & Services	53.804416	38.566343	11.544872
3	Finance & Technology	50.632184	23.821752	12.465465
4	Public Administration	39.056738	21.015625	15.015385
5	Trade & Transportation	32.657258	29.108696	12.584034
6	Manufacturing & Construction	20.328605	22.257143	10.690355

1.18 Problem 16

Write pandas code that replicates the output of the following SQL code:

```
SELECT gender_balance, premiums, COUNT(*)
FROM whpps
GROUP BY gender_balance, premiums
HAVING gender_balance is NOT NULL and premiums is NOT NULL;
```

[2 points]

```
[ ]: myquery16 = '''
SELECT gender_balance, premium_payment, COUNT(*)
FROM whpps
GROUP BY gender_balance, premium_payment
HAVING gender_balance is NOT NULL and premium_payment is NOT NULL
'''
pd.read_sql_query(myquery16, whadb)
```

```
[ ]: gender_balance premium_payment COUNT(*)
0      Balanced      Full      226
1      Balanced      No       77
2      Balanced      Partial  271
3      Mostly men      Full  293
4      Mostly men      No    87
5      Mostly men      Partial 321
6      Mostly women    Full  267
7      Mostly women    No   107
8      Mostly women    Partial 333
```

```
[38]: whpps16 = whpps.groupby(['gender_balance', 'premium_payment']).size().
      ↪reset_index(name='count').sort_values('gender_balance', ascending=True)
whpps16
```

```
[38]: gender_balance premium_payment count
0      Mostly men      Full    293
1      Mostly men      No     87
2      Mostly men      Partial 321
3      Balanced      Full    226
4      Balanced      No     77
```

5	Balanced	Partial	271
6	Mostly women	Full	267
7	Mostly women	No	107
8	Mostly women	Partial	333

The sort didn't work for some reason, but it's all there. Balanced is always in the middle no matter how I sort.